Tagging Products Using Image Classification

Brian Tomasik, Phyo Thiha and Douglas Turnbull
{btomasi1,pthiha1}@alum.swarthmore.edu, turnbull@cs.swarthmore.edu
Computer Science Department, Swarthmore College, Swarthmore, PA 19081

1. Introduction

Goal

- Develop a system to automatically **annotating products with labels**

Approach

- "Bag of visual words" image classifier
- Scale Invariant Feature Transform (SIFT)
- Hierarchical visual vocabulary
- Variant of nearest-neighbor classification

Tasks

- Classifying product images
  - Velcro vs. Lace shoes
  - Collar vs. V-Neck vs. Crew-Neck shirts
- Investigate the effect of
  - Numbers of product training examples
  - Multiple views of products in classification

2. Data Collection

- Shoe and shirt images
- ~ **3500 images** from Amazon.com and other online stores
- Labelled images with **category** and **viewpoint**

Table 1: The product categories collected. The vertical lines separate the classification tasks carried out in our experiments.

3. The *Bag of Visual Words* Approach

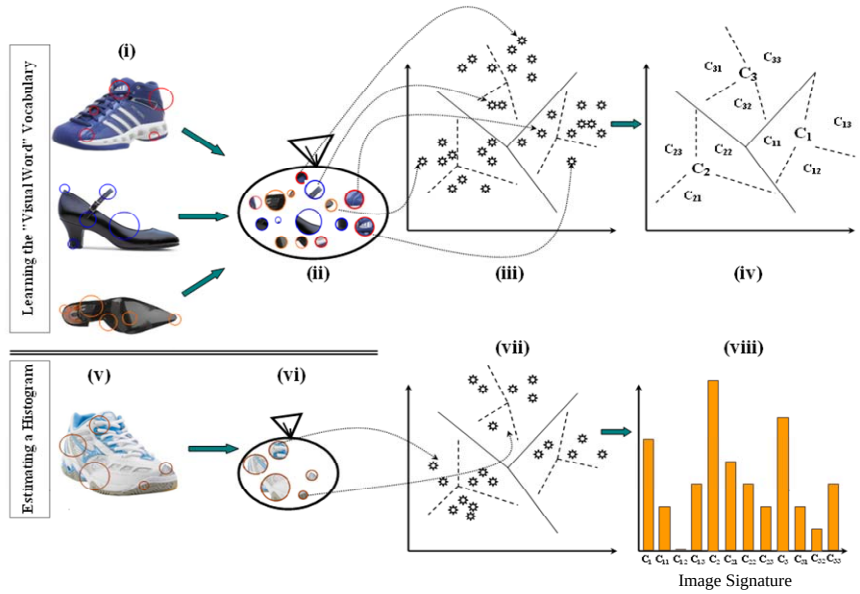

Figure 1: The first row shows the process of learning a vocabulary of **visual words** by **(i)** selecting **keypoints** from each image, **(ii) - (iii)** computing **SIFT descriptor vectors** at those keypoints, and **(iv) clustering** the entire collection of SIFT descriptors into groups whose centers will define the visual words. We **cluster into *k* groups** (k = 3 shown, k = 100 used) and then **recursively cluster** each of those groups to create a **tree of cluster centers**. The second row shows how we use the **visual-word tree**. **(v)** Given an image, we **(vi)** again compute SIFT descriptors at keypoints and then **(vii) walk each descriptor down the vocabulary tree** using **the closest cluster centers**. Each time a descriptor walks through a cluster center, we increment **the frequency count** for that visual word. **(viii)** The result is a **histogram of visual-word counts,** namely **image signatures**.

4 (i). Image Classification

- TF-IDF for Image Signatures, s_w
- Cosine-similarity for distance measures between s_{wi} and s_{wj}
- Variant of *k*-NN namely *Z-Score Voting* where weight, w_i , is:

$$w_i = -zscore(d(s_t, s_i)) = \frac{\mu_t - d(s_t, s_i)}{\sigma_t}$$

- Keypoints 10,000 with different selection methods: Canny edge detection, Random and combined

4 (ii). Multiple Views

- Each product has multiple associated images, corresponding to multiple views of the product
- Some viewpoints available are not helpful—e.g., underside for laced vs. velcro shoes
- Solution: Use all views available by calculating distances from each view of a product from all views of other products

5 (i). Training Set Size

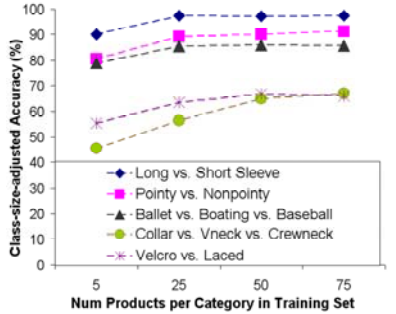

Figure 2. Class-size-adjusted accuracies saturated as we increased the number of images used in the training process for each product beyond 50.

5 (ii). Number of Views

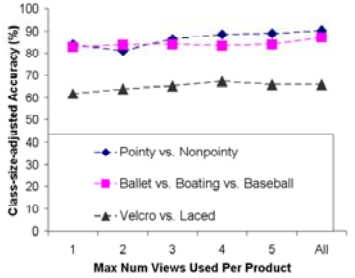

Figure 3. Class-size-adjusted accuracies improved as we increased the number of views used for some products.

References

- J. Sivic and A. Zisserman. Video google: A text retrieval approach to object matching in videos. In *ICCV '03*.
- D. Lowe. Distinctive image features from scale-invariant keypoints. *IJCV '04*, 60(2):91-110.
- D. Nister and H. Stewenius. Scalable recognition with a vocabulary tree. In *CVPR '06*, pages 2161-2168.
- A. Vedaldi. Bag of features. http://www.vlfeat.org/~vedaldi/code/bag/bag.html.
- B. Tomasik, P. Thiha, and D. Turnbull. Tagging products using image classification. Technical report, Swarthmore College, 2009. http://www.sccs.swarthmore.edu/users/09/btomasi1/tagging-products.html